

# Production Machine Learning Systems

Here are the assembled readings provided in this course.

## Module 1: Architecting Production ML Systems

- [Architecture of a real-world Machine Learning system](#)
- [Machine Learning Pipeline: Architecture of ML Platform in Production](#)
- [Design Decisions for Architecting Production Machine Learning Systems](#)
- [Production ML Systems](#)
- [3 Building Blocks of Machine Learning you Should Know as a Data Scientist](#)
- [MLOps: Continuous delivery and automation pipelines in machine learning](#)
- [Smart Decisions Game: Machine Learning for Architects](#)
- [Training and Serving CARET models using AI Platform Custom Containers and Cloud Run](#)
- [Using TensorFlow to predict product weight and dimensions](#)
- [Getting batch predictions](#)
- [How to extend a canned TensorFlow Estimator](#)
- [Introduction to loading data](#)
- [Google Cloud Vertex AI](#)
- [AI Simplified](#)
- [Cloud Blog AI & Machine Learning](#)
- [GitHub - Google Cloud Pipeline Components](#)
- [Vertex AI: Building a fraud detection model with AutoML](#)

## Module 2: Designing Adaptable ML Systems

- [Deep Learning AI Needs Tools To Adapt To Changes In The Data Environment](#)
- [MACHINE LEARNING FOR FUTURE SYSTEM DESIGNS](#)
- [Three Risks in Building Machine Learning Systems](#)
- [Advantages of Adaptive AI Over Traditional Machine Learning Models](#)
- [ML Opening New Doors For FPGAs](#)
- [Rules of Machine Learning: Best Practices for ML Engineering](#)
- [Best Practices for creating training data](#)
- [Productionizing Behavioural Features for Machine Learning with Apache Spark Streaming](#)
- [TensorFlow Data Validation: Checking and analyzing your data](#)

## Module 3: Designing High-performance ML Systems

- [How to Evaluate the Performance of Your Machine Learning Model](#)
- [Best practices for performance and cost optimization for machine learning](#)
- [How To Improve Machine Learning Model Performance: Five Ways](#)
- [Distributed TensorFlow model training on Cloud AI Platform \(TF Dev Summit '20\)](#)
- [Distributed training with TensorFlow](#)
- [Speeding Up Neural Network Training with Data Echoing](#)
- [Machine Learning Performance Improvement Cheat Sheet](#)
- [Building a High-Performance Data Pipeline with Tensorflow 2.x](#)



- [Distributed training with TensorFlow](#)
- [AutoML Tables](#)

## Module 4: Hybrid ML Systems

- [Kubeflow](#)
- [Introduction to Kubeflow](#)
- [Orchestrating TFX Pipelines](#)
- [Introduction to Machine Learning Pipelines with Kubeflow](#)
- [Kubeflow – a machine learning toolkit for Kubernetes](#)
- [ML for Mobile and Edge Devices - TensorFlow Lite](#)
- [TensorFlow Lite Examples | Machine Learning Mobile Apps](#)
- [Optimize TensorFlow models for mobile and embedded devices](#)
- [The Essential Guide To Learn TensorFlow Mobile and Tensorflow Lite](#)